

3.2 Regression

August 14, 2020

Table of Contents

- 1 Ideas
- 2 Libs and import data
- 3 Data Wrangle
 - 3.1 Split data train, test
 - 3.2 Fill NaN values
 - 3.3 Additional cleanup
 - 3.4 Split features (meteorological inputs) and label (PM2.5)
- 4 Linear Regression
 - 4.1 evaluate model performance
- 5 DecisionTree
 - 5.1 RandomForest
 - 5.2 Cross validation
 - 5.3 save model
 - 5.4 Grid Search
 - 5.5 Analyze model
 - 5.6 Scipy interval 95%
 - 5.7 Ensemble Methods
- 6 DarkSky Dataset
 - 6.1 Merge data
 - 6.2 Split train and test
 - 6.3 Train and validate
 - 6.3.1 Trainset
 - 6.3.2 Testset
- 7 RMSE

1 Ideas

- Load data from the last exercise (cleaned data including the best combination of MERRA-2 product and observed ground data)
- Let try with simple prediction using Linear Regression, Logistic Regression, Decision Tree and RandomForest Regression
- Measured the accurate (or the error) from those technique
- Apply the outcome with a forecast data source such as from DarkSky

2 Libs and import data

```
[1]: import warnings
warnings.filterwarnings('ignore')
```

```
[2]: import pandas as pd
```

```
[3]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
[4]: import numpy as np
```

```
[44]: # if you have not install sklearn (or scikit-learn) and scipy then install it
# ! pip install scikit-learn --user
# ! pip install scipy --user
```

```
[5]: # let pick up the data has been cleaned up the previous exercise, v3
df = pd.read_csv('data/comb_PM25_wind_Hanoi_2018_v3.csv',
                 parse_dates=['DATE'],
                 index_col=['DATE'])
df.head()
```

```
[5]:
```

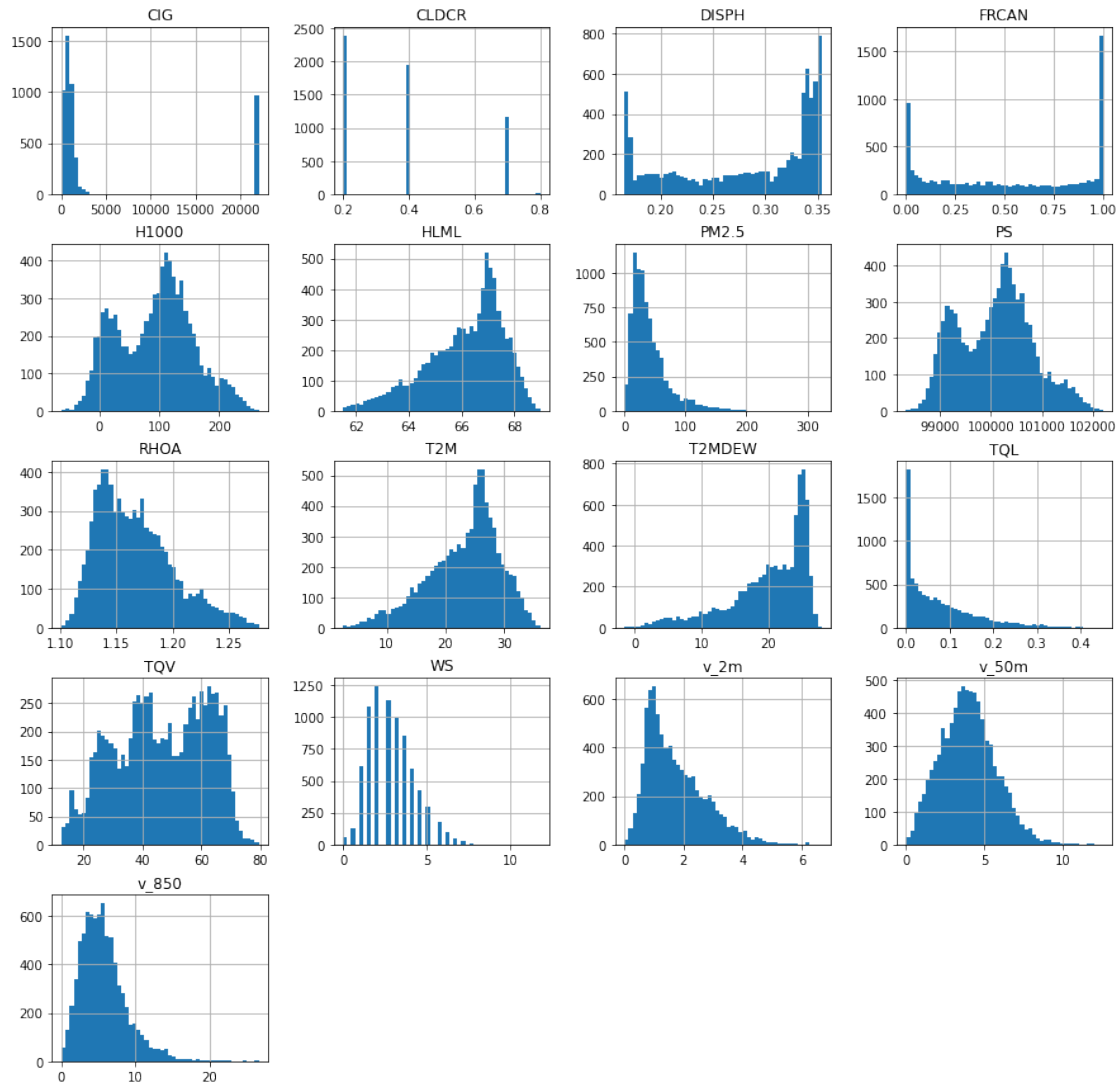
	PM2.5	T2MDEW	T2M	PS	TQV	\
DATE						
2018-01-01 01:00:00	69.2	10.79443	13.64376	100865.09	34.909637	
2018-01-01 02:00:00	75.5	10.72836	13.33932	100819.56	35.195385	
2018-01-01 03:00:00	90.2	10.60630	13.09753	100793.71	35.590984	
2018-01-01 04:00:00	97.6	10.50967	12.81360	100791.80	35.827934	
2018-01-01 05:00:00	89.1	10.49365	12.71010	100808.45	35.953880	

	TQL	H1000	DISPH	FRCAN	HLML	\
DATE						
2018-01-01 01:00:00	0.009235	160.25461	0.256226	1.000000	63.907425	
2018-01-01 02:00:00	0.006260	156.44829	0.256226	1.000000	63.832478	
2018-01-01 03:00:00	0.003489	154.54437	0.256104	0.993164	63.766266	
2018-01-01 04:00:00	0.002314	154.16837	0.255981	0.927490	63.718185	

2018-01-01 05:00:00 0.001480 155.72943 0.255859 0.786133 63.684280

	RHOA	CIG	WS	CLDCR	v_2m	v_50m	v_850
DATE							
2018-01-01 01:00:00	1.216159	NaN	NaN	NaN	0.429060	0.979452	6.175777
2018-01-01 02:00:00	1.217125	NaN	NaN	NaN	0.383400	0.872296	5.997708
2018-01-01 03:00:00	1.218085	NaN	NaN	NaN	0.339189	0.729687	5.915825
2018-01-01 04:00:00	1.218972	NaN	NaN	NaN	0.305853	0.666341	5.885087
2018-01-01 05:00:00	1.219831	NaN	NaN	NaN	0.258492	0.541109	5.796022

```
[6]: df.hist(bins=50, figsize=(15,15));
```



3 Data Wrangle

3.1 Split data train, test

```
[6]: # the idea is the split the dataset into two parts, one for training, and one
      ↪ for validation
def split_train_test(data, test_ratio):
    shuffled_indices = np.random.permutation(len(data))
    test_set_size = int(len(data)*test_ratio)
    test_indices = shuffled_indices[:test_set_size]
    train_indices = shuffled_indices[test_set_size:]
    return data.iloc[train_indices], data.iloc[test_indices]
```

```
[7]: # now we can split them with a ratio like this
train_set, test_set = split_train_test(df, 0.2)
```

```
[8]: len(train_set)
```

```
[8]: 6493
```

```
[9]: len(test_set)
```

```
[9]: 1623
```

```
[11]: # however, the sklearn library has such unility
from sklearn.model_selection import train_test_split
```

```
[13]: train_set, test_set = train_test_split(df, test_size=0.2, random_state=2020)
print(len(train_set))
train_set.head()
```

```
6492
```

```
[13]:
```

	PM2.5	T2MDEW	T2M	PS	TQV	\
DATE						
2018-11-01 01:00:00	66.0	13.76476	15.91640	100623.33	23.115633	
2018-03-11 13:00:00	97.1	16.36712	23.94467	100729.96	37.070614	
2018-11-02 17:00:00	28.0	15.21340	22.42794	100487.85	20.839620	
2018-03-27 23:00:00	63.7	18.28860	18.51727	100447.27	38.934628	
2018-12-01 04:00:00	55.0	20.01226	20.44995	100557.99	42.091070	

	TQL	H1000	DISPH	FRCAN	HLML	\
DATE						
2018-11-01 01:00:00	0.000000	140.67795	0.346069	0.000000	65.119470	
2018-03-11 13:00:00	0.095612	151.08076	0.170471	0.558838	65.698654	
2018-11-02 17:00:00	0.000000	130.45107	0.345093	0.000000	66.063030	
2018-03-27 23:00:00	0.004002	126.26744	0.192688	0.044815	65.261230	

```
2018-12-01 04:00:00 0.090057 136.40723 0.316406 0.208740 65.746920
```

```

                                RHOA      CIG  WS  CLDCR      v_2m      v_50m  \
DATE
2018-11-01 01:00:00 1.190721      NaN  2.1    NaN  1.317137  5.151577
2018-03-11 13:00:00 1.181534      NaN  2.1    0.4  3.250031  4.604924
2018-11-02 17:00:00 1.172179      NaN  1.0    NaN  0.614779  1.535728
2018-03-27 23:00:00 1.186127 22000.0  3.1    NaN  1.117351  4.444736
2018-12-01 04:00:00 1.178689   792.0  1.5    0.4  0.629109  1.737746

```

```

                                v_850
DATE
2018-11-01 01:00:00 13.778793
2018-03-11 13:00:00  5.120155
2018-11-02 17:00:00  3.374469
2018-03-27 23:00:00  6.082646
2018-12-01 04:00:00  5.265891

```

3.2 Fill NaN values

- Data are likely assembled with missing values
- Regression or machine learning works better with completed dataset

```
[14]: # make a copy and test out
df1 = df.copy(deep=True)
```

```
[15]: # calculate the median of the Windspeed (WS) input
median = df1['WS'].median()
```

```
[19]: df1.info()
```

```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8116 entries, 2018-01-01 01:00:00 to 2018-12-31 00:00:00
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PM2.5       8116 non-null    float64
1   T2MDEW      8116 non-null    float64
2   T2M         8116 non-null    float64
3   PS          8116 non-null    float64
4   TQV         8116 non-null    float64
5   TQL         8116 non-null    float64
6   H1000       8116 non-null    float64
7   DISPH       8116 non-null    float64
8   FRCAN       8092 non-null    float64
9   HLML        8092 non-null    float64
10  RHOA        8092 non-null    float64

```

```

11 CIG      5088 non-null   float64
12 WS       7809 non-null   float64
13 CLDCR    5502 non-null   float64
14 v_2m     8116 non-null   float64
15 v_50m    8116 non-null   float64
16 v_850    8116 non-null   float64
dtypes: float64(17)
memory usage: 1.1 MB

```

```
[20]: df1['WS'].fillna(median, inplace=True)
```

```
[21]: df1.info()
```

```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8116 entries, 2018-01-01 01:00:00 to 2018-12-31 00:00:00
Data columns (total 17 columns):
#   Column  Non-Null Count  Dtype
---  -
0   PM2.5    8116 non-null   float64
1   T2MDEW   8116 non-null   float64
2   T2M      8116 non-null   float64
3   PS       8116 non-null   float64
4   TQV      8116 non-null   float64
5   TQL      8116 non-null   float64
6   H1000    8116 non-null   float64
7   DISPH    8116 non-null   float64
8   FRCAN    8092 non-null   float64
9   HLML     8092 non-null   float64
10  RHOA     8092 non-null   float64
11  CIG      5088 non-null   float64
12  WS       8116 non-null   float64
13  CLDCR    5502 non-null   float64
14  v_2m     8116 non-null   float64
15  v_50m    8116 non-null   float64
16  v_850    8116 non-null   float64
dtypes: float64(17)
memory usage: 1.1 MB

```

```

[22]: # we can manually do this by for loop
for col in df1.columns:
    df1[col].fillna(df1[col].median(), inplace=True)
df1.info()

```

```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8116 entries, 2018-01-01 01:00:00 to 2018-12-31 00:00:00
Data columns (total 17 columns):
#   Column  Non-Null Count  Dtype
---  -

```

```

0    PM2.5    8116 non-null    float64
1    T2MDEW  8116 non-null    float64
2    T2M      8116 non-null    float64
3    PS       8116 non-null    float64
4    TQV      8116 non-null    float64
5    TQL      8116 non-null    float64
6    H1000    8116 non-null    float64
7    DISPH    8116 non-null    float64
8    FRCAN    8116 non-null    float64
9    HLML     8116 non-null    float64
10   RHOA     8116 non-null    float64
11   CIG      8116 non-null    float64
12   WS       8116 non-null    float64
13   CLDCR    8116 non-null    float64
14   v_2m     8116 non-null    float64
15   v_50m    8116 non-null    float64
16   v_850    8116 non-null    float64

```

dtypes: float64(17)

memory usage: 1.1 MB

```
[23]: # and sklearn has the Class to do such
      from sklearn.impute import SimpleImputer
```

```
[24]: inputer = SimpleImputer(strategy='median')
```

```
[25]: df2 =df.copy(deep=True)
      df2.info()
```

<class 'pandas.core.frame.DataFrame'>

DatetimeIndex: 8116 entries, 2018-01-01 01:00:00 to 2018-12-31 00:00:00

Data columns (total 17 columns):

#	Column	Non-Null Count	Dtype
0	PM2.5	8116 non-null	float64
1	T2MDEW	8116 non-null	float64
2	T2M	8116 non-null	float64
3	PS	8116 non-null	float64
4	TQV	8116 non-null	float64
5	TQL	8116 non-null	float64
6	H1000	8116 non-null	float64
7	DISPH	8116 non-null	float64
8	FRCAN	8092 non-null	float64
9	HLML	8092 non-null	float64
10	RHOA	8092 non-null	float64
11	CIG	5088 non-null	float64
12	WS	7809 non-null	float64
13	CLDCR	5502 non-null	float64
14	v_2m	8116 non-null	float64

```
15 v_50m 8116 non-null float64
16 v_850 8116 non-null float64
dtypes: float64(17)
memory usage: 1.1 MB
```

```
[26]: # evaluate df2
inputer.fit(df2)
```

```
[26]: SimpleImputer(add_indicator=False, copy=True, fill_value=None,
missing_values=nan, strategy='median', verbose=0)
```

```
[27]: # see the statistic, median in this case
inputer.statistics_
```

```
[27]: array([3.20000000e+01, 2.17670550e+01, 2.45143500e+01, 1.00156612e+05,
4.69396705e+01, 5.37109380e-02, 1.01305442e+02, 3.08166505e-01,
5.11840850e-01, 6.64740250e+01, 1.16165940e+00, 9.45000000e+02,
2.60000000e+00, 4.00000000e-01, 1.48621393e+00, 3.92495231e+00,
5.23608490e+00])
```

```
[28]: # transform is doing the work
df2full = inputer.transform(df1)
```

```
[29]: type(df2full)
```

```
[29]: numpy.ndarray
```

```
[30]: # convert the inputted dataset and and compared
df2 = pd.DataFrame(data=df2full, columns=df1.columns)
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8116 entries, 0 to 8115
Data columns (total 17 columns):
#   Column  Non-Null Count  Dtype
---  -
0   PM2.5    8116 non-null    float64
1   T2MDEW   8116 non-null    float64
2   T2M      8116 non-null    float64
3   PS       8116 non-null    float64
4   TQV      8116 non-null    float64
5   TQL      8116 non-null    float64
6   H1000    8116 non-null    float64
7   DISPH    8116 non-null    float64
8   FRCAN    8116 non-null    float64
9   HLML     8116 non-null    float64
10  RHOA     8116 non-null    float64
11  CIG      8116 non-null    float64
```



```

12 WS      8116 non-null float64
13 CLDCR   8116 non-null float64
14 v_2m    8116 non-null float64
15 v_50m   8116 non-null float64
16 v_850   8116 non-null float64
dtypes: float64(17)
memory usage: 1.1 MB

```

```

[32]: # you can save to a file with all missing values filled
df.to_csv('data/filled_PM2.5_Hanoi_2018.csv')

```

```


[31]: # let see if they are matched
fig, ax = plt.subplots(figsize=(15,6))
ax.plot(df1.WS.to_list())
ax.plot(df2.WS, 'ro')
ax.set_xlim(0,200)

```

```

[31]: (0.0, 200.0)

```



3.3 Additional cleanup

```


[32]: # let return back to the original dataset (df) before fill up NaN
sns.heatmap(df.corr(), cmap='seismic')

```

```

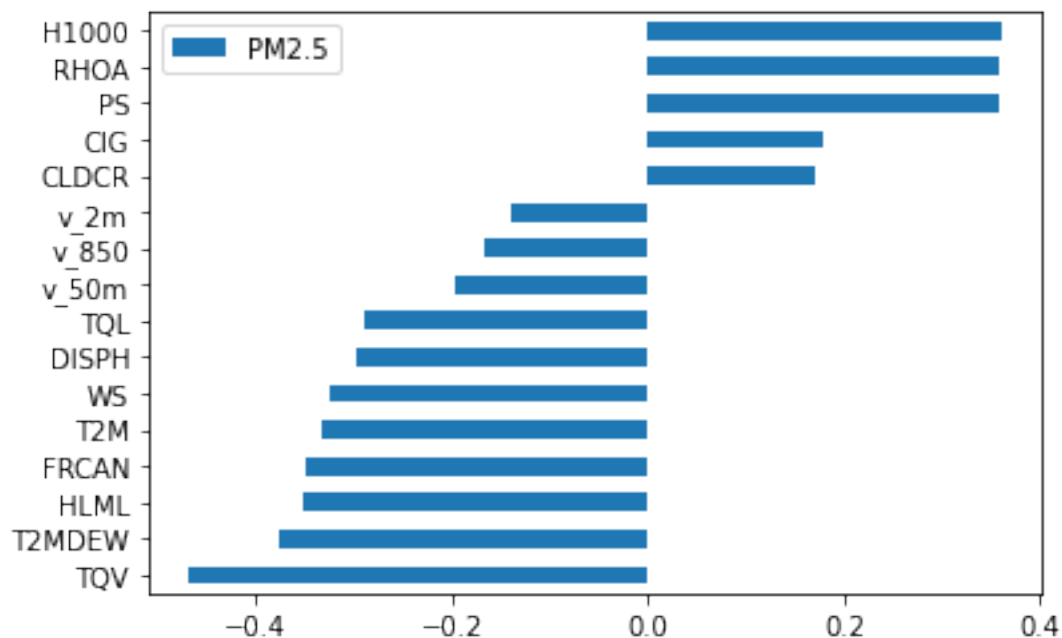
[32]: <AxesSubplot:>

```



```
[33]: # and only correlation with PM2.5
df.corr()['PM2.5'].sort_values().to_frame().drop('PM2.5').plot.barh()
```

[33]: <AxesSubplot:>



```
[34]: df.columns
```

```
[34]: Index(['PM2.5', 'T2MDEW', 'T2M', 'PS', 'TQV', 'TQL', 'H1000', 'DISPH', 'FRCAN',  
        'HLML', 'RHOA', 'CIG', 'WS', 'CLDCR', 'v_2m', 'v_50m', 'v_850'],  
        dtype='object')
```

```
[35]: # drop some columns either weak in correlation or dependent (redundant) to other  
      ↪ inputs  
df.drop(columns=['CLDCR', 'v_2m', 'v_50m', 'v_850', 'FRCAN', 'DISPH'],  
      ↪ inplace=True)
```

```
[36]: df.head()
```

```
[36]:
```

		PM2.5	T2MDEW	T2M	PS	TQV	\
DATE							
2018-01-01 01:00:00		69.2	10.79443	13.64376	100865.09	34.909637	
2018-01-01 02:00:00		75.5	10.72836	13.33932	100819.56	35.195385	
2018-01-01 03:00:00		90.2	10.60630	13.09753	100793.71	35.590984	
2018-01-01 04:00:00		97.6	10.50967	12.81360	100791.80	35.827934	
2018-01-01 05:00:00		89.1	10.49365	12.71010	100808.45	35.953880	

		TQL	H1000	HLML	RHOA	CIG	WS
DATE							
2018-01-01 01:00:00		0.009235	160.25461	63.907425	1.216159	NaN	NaN
2018-01-01 02:00:00		0.006260	156.44829	63.832478	1.217125	NaN	NaN
2018-01-01 03:00:00		0.003489	154.54437	63.766266	1.218085	NaN	NaN
2018-01-01 04:00:00		0.002314	154.16837	63.718185	1.218972	NaN	NaN
2018-01-01 05:00:00		0.001480	155.72943	63.684280	1.219831	NaN	NaN

```
[64]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
DatetimeIndex: 8116 entries, 2018-01-01 01:00:00 to 2018-12-31 00:00:00  
Data columns (total 11 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   PM2.5        8116 non-null   float64  
1   T2MDEW       8116 non-null   float64  
2   T2M          8116 non-null   float64  
3   PS           8116 non-null   float64  
4   TQV          8116 non-null   float64  
5   TQL          8116 non-null   float64  
6   H1000        8116 non-null   float64  
7   HLML         8116 non-null   float64  
8   RHOA         8116 non-null   float64
```

```

9    CIG      8116 non-null    float64
10   WS       8116 non-null    float64
dtypes: float64(11)
memory usage: 760.9 KB

```

3.4 Split features (meteorological inputs) and label (PM2.5)

```
[59]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8116 entries, 2018-01-01 01:00:00 to 2018-12-31 00:00:00
Data columns (total 11 columns):
 #   Column      Non-Null Count  Dtype
---  -
0    PM2.5      8116 non-null   float64
1    T2MDEW     8116 non-null   float64
2    T2M        8116 non-null   float64
3    PS         8116 non-null   float64
4    TQV        8116 non-null   float64
5    TQL        8116 non-null   float64
6    H1000      8116 non-null   float64
7    HLML       8092 non-null   float64
8    RHOA       8092 non-null   float64
9    CIG        5088 non-null   float64
10   WS         7809 non-null   float64
dtypes: float64(11)
memory usage: 760.9 KB

```

```
[41]: # let make X as the matrix for the feature (or inputs)
X = df.drop('PM2.5', axis=1)
```

```
[37]: # and lowercase y as the label (or the value of the target)
y = df['PM2.5'].copy()
```

```
[53]: # let build the input instance to work with whole data at one
# to input more than one columns, we can use this
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

```
[58]: # transform data from real value to a relative to the set
from sklearn.preprocessing import StandardScaler
```

```
[54]: # we have all column with numeric values
num_attrs = list(df.columns)
num_attrs.remove('PM2.5')
num_attrs
```

```
[54]: ['T2MDEW', 'T2M', 'PS', 'TQV', 'TQL', 'H1000', 'HLML', 'RHOA', 'CIG', 'WS']
```

```
[56]: # first is the strategy for imputer using median
# then convert the absolute value in the each column using the Standard Class
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='median')),
    ('std_scaler', StandardScaler()),
])
```

```
[57]: # and instance to transform all column at one
full_pipeline = ColumnTransformer([
    ('num', num_pipeline, num_attrs)
])
```

```
[60]: # or building a function to do all in one step
def impute_transform(data=None):
    num_pipeline = Pipeline([
        ('imputer', SimpleImputer(strategy='median')),
        ('std_scaler', StandardScaler()),
    ])
    num_attrs = list(data.columns)
    full_pipeline = ColumnTransformer([
        ('num', num_pipeline, num_attrs)
    ])
    return full_pipeline.fit_transform(data) # return a numpy array
```

```
[61]: X_scaled = impute_transform(data=X)
```

```
[66]: # how do we know that the data has been fixed properly
X_scaled_test = impute_transform(data=X)
dft = pd.DataFrame(data=X_scaled_test, columns=num_attrs)
dft.info()
# looking good
del dft
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 8116 entries, 0 to 8115
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	T2MDEW	8116 non-null	float64
1	T2M	8116 non-null	float64
2	PS	8116 non-null	float64
3	TQV	8116 non-null	float64
4	TQL	8116 non-null	float64
5	H1000	8116 non-null	float64
6	HLML	8116 non-null	float64

```
7  RHOA      8116 non-null   float64
8  CIG       8116 non-null   float64
9  WS        8116 non-null   float64
dtypes: float64(10)
memory usage: 634.2 KB
```

```
[67]: # now we can split data, the test_size is the portion of data use for validation
      # random_state is to provide consistent if you want to replicate the result
      X_train, X_test, y_train, y_test = train_test_split(X_scaled,y, test_size=0.33,
      ↪random_state=2020)
```

```
[68]: len(X_train), len(X_test)
```

```
[68]: (5437, 2679)
```

4 Linear Regression

```
[43]: # Linear regression is the simplest form of data learning, let try this first
      from sklearn.linear_model import LinearRegression
```

```
[45]: # make an instant of the class
      lin_reg = LinearRegression()
```

```
[69]: # and training the model using the _train dataset
      lin_reg.fit(X_train, y_train)
```

```
[69]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
[70]: lin_reg.get_params()
```

```
[70]: {'copy_X': True, 'fit_intercept': True, 'n_jobs': None, 'normalize': False}
```

```
[71]: # let see the output of the mode
      lin_reg.coef_
```


```
[71]: array([ -3.15806877,   9.23068653, -237.20129647, -11.37237178,
          -1.51486159,  216.26310513,   19.10913327,   47.2122002 ,
           2.99459991,  -8.55437448])
```

```
[72]: # predict based on the training set
      y_train_prd = lin_reg.predict(X_train)
```

```
[73]: # let see how the label (y_train), and predicting of the label (y_prd)
      plt.figure(figsize=(15,5))
      plt.plot(y_train.to_list())
```

```
plt.plot(y_train_prd)
plt.xlim(0,500)
```


[73]: (0.0, 500.0)



```
[74]: # more important, how about with validation set (test set)
y_test_prd = lin_reg.predict(X_test)
```

```
[75]: plt.figure(figsize=(15,5))
plt.plot(y_test.to_list())
plt.plot(y_test_prd)
plt.xlim(0,500)
```

[75]: (0.0, 500.0)



- kind of there, but not easy to assess

4.1 evaluate model performance

```
[76]: # for numeric data, one simple way to see how far
      # between the prediction and the target
      from sklearn.metrics import mean_squared_error
```

```
[77]: # on training set
      lin_train_mse = mean_squared_error(y_train, y_train_prd)
      print('Trainset: Root Squared Mean Error', np.sqrt(lin_train_mse))
```

Trainset: Root Squared Mean Error 25.50502509220003

```
[78]: # on test set
      lin_test_mse = mean_squared_error(y_test, y_test_prd)
      print('Test set: Root Squared Mean Error', np.sqrt(lin_test_mse))
```

Test set: Root Squared Mean Error 25.924922211091516

```
[79]: # the average value label set (y set)
      y.mean()
```

[79]: 40.75873583045832

```
[80]: # relative error
      print(f'Relative Error: {100*np.sqrt(lin_test_mse)/y.mean():.0f}%')
```

Relative Error: 64%

- underfitting (data), we have can predict the PM2.5 concentration, but bear in mind the value can be off 64% (up and down), and that only works 66 out of 100 chances (given the error is random)
- let try another model, also, make a dictionary to keep the score between model

```
[95]: results = dict()
      def add_stats(model=None, train_rmse=None, test_rmse=None):
          global results
          results[model] = {'train_rmse': round(train_rmse,1),
                           'test_rmse': round(test_rmse, 1)}
          return None
```

```
[97]: add_stats(model='linear reg',
               train_rmse=np.sqrt(lin_train_mse),
               test_rmse=np.sqrt(lin_test_mse))
      results
```

[97]: {'linear reg': {'train_rmse': 25.5, 'test_rmse': 25.9}}

5 DecisionTree

```
[98]: # let try the same approach with decision tree
      from sklearn.tree import DecisionTreeRegressor

[99]: tree_reg = DecisionTreeRegressor()

[100]: tree_reg.fit(X_train, y_train)

[100]: DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
                             max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort='deprecated',
                             random_state=None, splitter='best')

[101]: tree_reg.feature_importances_

[101]: array([0.28720845, 0.06360633, 0.05308155, 0.14258232, 0.10199783,
              0.06372484, 0.05821381, 0.04296177, 0.07141657, 0.11520652])

[102]: tree_ytrain_d = tree_reg.predict(X_train)

[104]: # no error, too good!
      tree_train_rmse = np.sqrt(mean_squared_error(y_train, tree_ytrain_d))
      tree_train_rmse

[104]: 0.0

[105]: tree_ytest_d = tree_reg.predict(X_test)

[110]: tree_test_rmse = np.sqrt(mean_squared_error(y_test, tree_ytest_d))
      tree_test_rmse # higher than regression,

[110]: 28.866692614774283


[111]: # let bag the result
      add_stats(model='decisiontree reg',
                 train_rmse=tree_train_rmse,
                 test_rmse=tree_test_rmse)
      results

[111]: {'linear reg': {'train_rmse': 25.5, 'test_rmse': 25.9},
       'decisiontree reg': {'train_rmse': 0.0, 'test_rmse': 28.9}}


[107]: # the results by train set and test set are rather different, to see it
      def plot_prediction(label=None, prediction=None, limit=200):
```

```
plt.figure(figsize=(14,6))
plt.plot(label.to_list())
plt.plot(prediction, 'ro')
plt.xlim(0, limit)
return None
```

```
[108]: plot_prediction(y_train, tree_ytrain_d)
```



```
[109]: plot_prediction(y_test, tree_ytest_d)
```



- overfitting data for training set (memorization), and underfitting with test set (not map key feature)

5.1 RandomForest

```
[112]: # more powerful model
from sklearn.ensemble import RandomForestRegressor
```

```
[113]: forest_reg = RandomForestRegressor()
forest_reg.fit(X_train, y_train)
```

```
[113]: RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             max_samples=None, min_impurity_decrease=0.0,
                             min_impurity_split=None, min_samples_leaf=1,
                             min_samples_split=2, min_weight_fraction_leaf=0.0,
                             n_estimators=100, n_jobs=None, oob_score=False,
                             random_state=None, verbose=0, warm_start=False)
```

```
[114]: forest_ytrain_p = forest_reg.predict(X_train)
```

```
[116]: mse_train = mean_squared_error(y_train, forest_ytrain_p)
rmse_train = np.sqrt(mse_train)
rmse_train
```

```
[116]: 7.319539455569634
```

```
[117]: # test set
forest_ytest_p = forest_reg.predict(X_test)
```

```
[119]: mse_test = mean_squared_error(y_test, forest_ytest_p)
rmse_test = np.sqrt(mse_test)
rmse_test
```

```
[119]: 19.94340670771594
```

```
[131]: add_stats(model='randomforest reg',
              train_rmse=rmse_train,
              test_rmse=rmse_test)
results
```

```
[131]: {'linear reg': {'train_rmse': 25.5, 'test_rmse': 25.9},
      'decisiontree reg': {'train_rmse': 0.0, 'test_rmse': 28.9},
      'randomforest reg': {'train_rmse': 7.3, 'test_rmse': 19.9}}
```

5.2 Cross validation

```
[122]: from sklearn.model_selection import cross_val_score
```

```
[124]: scores = cross_val_score(tree_reg, X_train, y_train,  
                                scoring='neg_mean_squared_error', cv=10)
```

```
[125]: tree_rmse_scores = np.sqrt(-scores)  
tree_rmse_scores
```

```
[125]: array([25.0286516 , 27.29367756, 25.45272398, 28.72647087, 27.21742825,  
            27.55288765, 25.1130188 , 31.19758898, 27.76135077, 28.17651903])
```

```
[126]: def display_scores(scores):  
        print("Scores: ", scores)  
        print("Mean: ", scores.mean())  
        print("Standard Deviation: ", scores.std())
```

```
[127]: display_scores(scores)
```

```
Scores: [-626.43340074 -744.94483456 -647.84115809 -825.21012868 -740.78840074  
        -759.16161765 -630.66371324 -973.28955801 -770.69259669 -793.91622468]  
Mean: -751.2941633050319  
Standard Deviation: 98.99266952818422
```

```
[128]: lin_scores = cross_val_score(lin_reg, X_train, y_train,  
                                    scoring='neg_mean_squared_error', cv=10)
```

```
[129]: display_scores(lin_scores)
```

```
Scores: [-550.85870613 -723.57425609 -604.74504505 -774.75436464 -558.6424609  
        -650.90144343 -694.9417229 -791.12250196 -464.54518714 -719.45029524]  
Mean: -653.3535983471008  
Standard Deviation: 101.11405749157925
```

```
[130]: forest_scores = cross_val_score(forest_reg, X_train, y_train,  
                                        scoring='neg_mean_squared_error', cv=10)
```

```
[132]: display_scores(forest_scores)
```

```
Scores: [-336.85249494 -414.55027717 -340.52606942 -495.02698419 -328.95593693  
        -347.62383281 -410.42979521 -546.46635029 -250.81217664 -405.05053205]  
Mean: -387.6294449641497  
Standard Deviation: 81.77438259082012
```

```
[133]: # how about on test set:  
for model in [lin_reg, tree_reg, forest_reg]:
```

```
scores = cross_val_score(model, X_test, y_test,
                          scoring='neg_mean_squared_error', cv=10)
display_scores(scores)
print('-'*40)
```

```
Scores: [-771.63467574 -554.4382247 -813.16490696 -717.49508355 -598.56766015
        -632.13746509 -569.73439949 -753.33939901 -661.54505262 -680.24197289]
Mean:    -675.229884018897
Standard Deviation:  83.5905310705698
-----
```

```
Scores: [-1259.97597015 -782.79813433 -1032.93242537 -1106.06947761
        -696.41026119 -846.71276119 -801.66145522 -820.10895522
        -866.49783582 -906.43479401]
Mean:    -911.9602070126894
Standard Deviation:  162.3358653447053
-----
```

```
Scores: [-602.63798754 -381.18848133 -554.00563056 -519.45501511 -321.74875161
        -450.15756396 -336.19587503 -490.46458475 -539.81659685 -420.10498245]
Mean:    -461.57754692068863
Standard Deviation:  90.73701388730139
-----
```

- look like the RandomForest Regression performs a bit better than the first two

5.3 save model

```
[135]: # just in case you want to save your work
from sklearn.externals import joblib
```

```
[134]: # import os
# os.makedirs('model')
```

```
[135]: joblib.dump(forest_reg, 'model/forest_reg.pkl' )
```

```
[135]: ['model/forest_reg.pkl']
```

5.4 Grid Search

```
[136]: # we want to model performs better, in this case we tune the hyperparameters
from sklearn.model_selection import GridSearchCV
```

```
[137]: param_grid = [{'n_estimators': [3, 10, 30], 'max_features': [2, 4, 6, 8]},
                     {'bootstrap': [False], 'n_estimators': [3, 10], 'max_features': 4},
                     {'bootstrap': [True], 'n_estimators': [3, 10], 'max_features': 4}],
```

```
[138]: forest_reg = RandomForestRegressor()
grid_search = GridSearchCV(forest_reg, param_grid, cv=5,
                           scoring='neg_mean_squared_error',
                           return_train_score=True)
```

```
[139]: grid_search.fit(X_train, y_train)
```

```
[139]: GridSearchCV(cv=5, error_score=nan,
                  estimator=RandomForestRegressor(bootstrap=True, ccp_alpha=0.0,
                                                    criterion='mse', max_depth=None,
                                                    max_features='auto',
                                                    max_leaf_nodes=None,
                                                    max_samples=None,
                                                    min_impurity_decrease=0.0,
                                                    min_impurity_split=None,
                                                    min_samples_leaf=1,
                                                    min_samples_split=2,
                                                    min_weight_fraction_leaf=0.0,
                                                    n_estimators=100, n_jobs=None,
                                                    oob_score=False, random_state=None,
                                                    verbose=0, warm_start=False),
                  iid='deprecated', n_jobs=None,
                  param_grid=[{'max_features': [2, 4, 6, 8],
                                'n_estimators': [3, 10, 30]},
                                {'bootstrap': [False], 'max_features': [2, 3, 4],
                                 'n_estimators': [3, 10]}],
                  pre_dispatch='2*n_jobs', refit=True, return_train_score=True,
                  scoring='neg_mean_squared_error', verbose=0)
```

```
[140]: # and see the best estimator
grid_search.best_estimator_
```

```
[140]: RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
                             max_depth=None, max_features=4, max_leaf_nodes=None,
                             max_samples=None, min_impurity_decrease=0.0,
                             min_impurity_split=None, min_samples_leaf=1,
                             min_samples_split=2, min_weight_fraction_leaf=0.0,
                             n_estimators=30, n_jobs=None, oob_score=False,
                             random_state=None, verbose=0, warm_start=False)
```

```
[141]: # or best parameters
grid_search.best_params_
```

```
[141]: {'max_features': 4, 'n_estimators': 30}
```

```
[142]: # or see the how each combination has worked
cvres = grid_search.cv_results_
```

```
[144]: for mean_score, params in zip(cvres['mean_test_score'], cvres['params']):
        print(round(np.sqrt(-mean_score),2), params)
```

```
24.44 {'max_features': 2, 'n_estimators': 3}
21.28 {'max_features': 2, 'n_estimators': 10}
20.33 {'max_features': 2, 'n_estimators': 30}
23.87 {'max_features': 4, 'n_estimators': 3}
21.03 {'max_features': 4, 'n_estimators': 10}
19.98 {'max_features': 4, 'n_estimators': 30}
23.36 {'max_features': 6, 'n_estimators': 3}
20.92 {'max_features': 6, 'n_estimators': 10}
20.0 {'max_features': 6, 'n_estimators': 30}
23.39 {'max_features': 8, 'n_estimators': 3}
21.0 {'max_features': 8, 'n_estimators': 10}
20.19 {'max_features': 8, 'n_estimators': 30}
23.2 {'bootstrap': False, 'max_features': 2, 'n_estimators': 3}
20.45 {'bootstrap': False, 'max_features': 2, 'n_estimators': 10}
22.49 {'bootstrap': False, 'max_features': 3, 'n_estimators': 3}
20.33 {'bootstrap': False, 'max_features': 3, 'n_estimators': 10}
22.55 {'bootstrap': False, 'max_features': 4, 'n_estimators': 3}
20.51 {'bootstrap': False, 'max_features': 4, 'n_estimators': 10}
```

- so the best case with Random Forest is 19.98 (g/m3) as the standard deviation

5.5 Analyze model

```
[145]: # could look back to see how the weight of each input
feature_importances = grid_search.best_estimator_.feature_importances_
feature_importances
```

```
[145]: array([0.16949425, 0.06773473, 0.09429158, 0.1824907 , 0.09555117,
            0.06977307, 0.0676954 , 0.07244233, 0.07794705, 0.10257973])
```

```
[147]: sorted(zip(feature_importances, X.columns), reverse=True)
```

```
[147]: [(0.18249070117732566, 'TQV'),
        (0.16949424694791368, 'T2MDEW'),
        (0.10257973245798097, 'WS'),
        (0.09555117432888298, 'TQL'),
        (0.09429157739434971, 'PS'),
        (0.07794704937895328, 'CIG'),
        (0.07244232993501565, 'RHOA'),
        (0.06977306516481928, 'H1000'),
        (0.0677347267676391, 'T2M'),
        (0.06769539644711989, 'HLML')]
```

- so the total liquid volumn and dewpoint influences PM2.5 more than the air temperature or surface level height

```
[152]: # let see how grid search performs
# train set
grid_ytrain_p = grid_search.predict(X_train)
grid_mse = mean_squared_error(y_train, grid_ytrain_p)
grid_train_rmse = np.sqrt(grid_mse)
grid_train_rmse
```

```
[152]: 7.795150104149395
```

```
[153]: # let see how the prediction look like after hypertunning
# test set
grid_ytest_p = grid_search.predict(X_test)
grid_test_mse = mean_squared_error(y_test, grid_ytest_p)
grid_test_rmse = np.sqrt(grid_test_mse)
grid_test_rmse
```


```
[153]: 20.033245999333058
```

```
[154]: # still more to work with, but let bag the result for later comparison
add_stats(model='gridsearch',
          train_rmse=grid_train_rmse,
          test_rmse=grid_test_rmse)
results
```

```
[154]: {'linear reg': {'train_rmse': 25.5, 'test_rmse': 25.9},
      'decisiontree reg': {'train_rmse': 0.0, 'test_rmse': 28.9},
      'randomforest reg': {'train_rmse': 7.3, 'test_rmse': 19.9},
      'gridsearch': {'train_rmse': 7.8, 'test_rmse': 20.0}}
```

```
[157]: # let visualize it
std_ = grid_test_rmse
xindex = np.arange(0, len(y_test))
fig, ax = plt.subplots(figsize=(15,6))
ax.plot(y_test.to_list(), color='black', label='measured')
ax.plot(grid_ytest_p, color='red', alpha=0.8, label='predicted')
ax.fill_between(xindex, grid_ytest_p-std_,grid_ytest_p+std_,
               color='red', alpha=0.3, label='1sigma' )
ax.fill_between(xindex, grid_ytest_p-2*std_,grid_ytest_p+2*std_,
               color='red', alpha=0.2, label = '2sigma')
ax.set_xlim(0,200)
ax.legend()
```

```
[157]: <matplotlib.legend.Legend at 0x7f43b3248ef0>
```

- an easy way to make sure the prediction range captures the real value is to increase the band, and so enlarge the uncertainty
- but let see how confidence we have on the RSME,

5.6 Scipy interval 95%

```
[158]: # let use stats from scipy library
from scipy import stats
```

```
[161]: # and look confidence of .95, or the area that a value will be inside the range
        ↪ with 95 chances of 100
confidence = 0.95
```

```
[162]: squared_errors = (grid_ytest_p - y_test)**2
```

```
[163]: np.sqrt(stats.t.interval(
        confidence,
        len(squared_errors)-1,
        loc=squared_errors.mean(),
        scale=stats.sem(squared_errors)))
```

```
[163]: array([18.83035947, 21.1678873 ])
```

- so we are pretty sure that standard deviation from grid search is from 18.8 to 21.1
- how confidence: 95 chances out of 100, this RMSE will be within this range

5.7 Ensemble Methods

```
[164]: # let look at a final approach to combine three regression we have so far using  
       ↳ Voting method  
       from sklearn.ensemble import RandomForestRegressor, VotingRegressor  
       from sklearn.linear_model import LinearRegression  
  
[165]: from sklearn.tree import DecisionTreeRegressor  
  
[166]: # re-define an instance, all training in the previous sessions are gone  
       lin_reg = LinearRegression()  
       tree_reg = DecisionTreeRegressor()  
       rnd_reg = RandomForestRegressor()  
  
[167]: # and make each model as an parameter for then ensemble (voting)  
       voting_reg = VotingRegressor(  
           estimators=[('lin', lin_reg),  
                       ('rnd', rnd_reg),  
                       ('tree', tree_reg)  
                       ],  
       )  
  
[168]: # train model by the train set  
       voting_reg.fit(X_train, y_train)  
  
[168]: VotingRegressor(estimators=[('lin',  
                                   LinearRegression(copy_X=True, fit_intercept=True,  
                                                    n_jobs=None, normalize=False)),  
                                   ('rnd',  
                                   RandomForestRegressor(bootstrap=True,  
                                                         ccp_alpha=0.0,  
                                                         criterion='mse',  
                                                         max_depth=None,  
                                                         max_features='auto',  
                                                         max_leaf_nodes=None,  
                                                         max_samples=None,  
                                                         min_impurity_decrease=0.0,  
                                                         min_impurity_split=None,  
                                                         min_samples_leaf=1,  
                                                         min_samples_split=2,  
                                                         mi...  
                                                         random_state=None, verbose=0,  
                                                         warm_start=False)),  
                                   ('tree',  
                                   DecisionTreeRegressor(ccp_alpha=0.0,  
                                                         criterion='mse',  
                                                         max_depth=None,
```

```

max_features=None,
max_leaf_nodes=None,
min_impurity_decrease=0.0,
min_impurity_split=None,
min_samples_leaf=1,
min_samples_split=2,
min_weight_fraction_leaf=0.0,
presort='deprecated',
random_state=None,
splitter='best'))],

n_jobs=None, weights=None)

```

```

[169]: # check estimator (parameters)
voting_reg.estimators_

```

```

[169]: [LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False),
RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
max_depth=None, max_features='auto', max_leaf_nodes=None,
max_samples=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=100, n_jobs=None, oob_score=False,
random_state=None, verbose=0, warm_start=False),
DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')]

```

```

[170]: # now we do predicting on the test set
y_entest_p = voting_reg.predict(X_test)

```

```

[171]: mse = mean_squared_error(y_entest_p, y_test)
en_test = np.sqrt(mse)

```

```

[172]: # we could run on train set
y_entrain_p = voting_reg.predict(X_train)
mse = mean_squared_error(y_entrain_p, y_train)
en_train = np.sqrt(mse)

```

```

[173]: en_test, en_train

```

```

[173]: (21.89413644598667, 10.545304085343597)


```

```
[174]: # still in 21 for test set
add_stats(model='voting reg',
          train_rmse=en_train,
          test_rmse=en_test)
results
```

```
[174]: {'linear reg': {'train_rmse': 25.5, 'test_rmse': 25.9},
       'decisiontree reg': {'train_rmse': 0.0, 'test_rmse': 28.9},
       'randomforest reg': {'train_rmse': 7.3, 'test_rmse': 19.9},
       'gridsearch': {'train_rmse': 7.8, 'test_rmse': 20.0},
       'voting reg': {'train_rmse': 10.5, 'test_rmse': 21.9}}
```

```
[175]: # let visualize the data from Ensemble with test set
plt.figure(figsize=(15,5))
plt.plot(y_test.to_list())
plt.plot(y_ensemble_p, lw=1)
plt.xlim(200,300)
```

```
[175]: (200.0, 300.0)
```



- so we can get a $RMSE = 20$ from several model, which is about 50% as the relative standard deviation
- this dataset is combined from several source, but not easy to get from a forecast product (in fact, I am struggling to get those), so we will try out a dataset with lessure feasture,

6 DarkSky Dataset

- you can check out this API at [DarkSky.net](https://darksky.net)

6.1 Merge data

```
[176]: # load data in
dk = pd.read_csv('data/darksky_hanoi_2018.csv', parse_dates=['time'],
↳index_col=['time'])
```

```
[177]: dk.columns
```

```
[177]: Index(['apparenttemperature', 'cloudcover', 'dewpoint', 'humidity', 'icon',
'ozone', 'precipintensity', 'precipprobability', 'preciptype',
'pressure', 'summary', 'temperature', 'uvindex', 'visibility',
'windbearing', 'windgust', 'windspeed'],
dtype='object')
```

```
[178]: # select few important columns
cols = ['temperature', 'dewpoint', 'humidity', 'pressure',
↳'precipintensity', 'cloudcover', 'visibility', 'windspeed']
```

```
[179]: dkt = dk[cols]
```

```
[180]: # load PM2.5 data
pm = pd.read_csv('data/cleaned_pm25_Hanoi_PM2.5_2018_YTD.csv',
parse_dates=['Date (LT)'],
index_col=['Date (LT)'])
```

```
[181]: pm.head()
```

```
[181]:
```

	pm25
Date (LT)	
2018-01-01 01:00:00	69.2
2018-01-01 02:00:00	75.5
2018-01-01 03:00:00	90.2
2018-01-01 04:00:00	97.6
2018-01-01 05:00:00	89.1

```
[182]: # check duplicated data if you want, wait, this is too much
dkt.duplicated().sum()
```

```
[182]: 1783
```

```
[183]: dkt.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8745 entries, 2017-12-31 00:00:00 to 2019-01-01 00:00:00
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
#   Column                Non-Null Count  Dtype
```

```

0   temperature      8742 non-null   float64
1   dewpoint         8742 non-null   float64
2   humidity         8742 non-null   float64
3   pressure         1137 non-null   float64
4   precipintensity  3383 non-null   float64
5   cloudcover       6862 non-null   float64
6   visibility       8711 non-null   float64
7   windspeed       8600 non-null   float64
dtypes: float64(8)
memory usage: 614.9 KB

```

```
[184]: # let sort index (datetime) first
      dct.sort_index(inplace=True)
```

```
[185]: # and see duplicated row
      dct.loc[dct.duplicated()].head(6)
```

```
[185]:
```

	temperature	dewpoint	humidity	pressure	\
time					
2017-12-31 05:00:00	16.98	9.99	0.63	NaN	
2017-12-31 06:00:00	16.98	9.99	0.63	NaN	
2018-01-05 06:00:00	19.99	19.00	0.94	NaN	
2018-01-05 08:00:00	19.99	19.00	0.94	NaN	
2018-01-05 09:00:00	19.99	19.00	0.94	NaN	
2018-01-06 23:00:00	19.99	19.00	0.94	NaN	

	precipintensity	cloudcover	visibility	windspeed
time				
2017-12-31 05:00:00	0.0	NaN	10.01	3.6
2017-12-31 06:00:00	0.0	NaN	10.01	3.6
2018-01-05 06:00:00	0.0	0.75	2.09	1.5
2018-01-05 08:00:00	0.0	0.75	1.50	1.5
2018-01-05 09:00:00	0.0	0.75	1.40	1.5
2018-01-06 23:00:00	0.0	0.75	3.51	2.6

- they are matched exact, but dropping them will need to fill in more data later, so it is okay to keep the closest values in by adjacent rows

```
[186]: # merge data
      df = pd.merge(pm, dct, right_index=True, left_index=True, how='left')
```

```
[187]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 8190 entries, 2018-01-01 01:00:00 to 2019-01-01 00:00:00
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -

```

```

0    pm25          8190 non-null    float64
1    temperature  8150 non-null    float64
2    dewpoint     8150 non-null    float64
3    humidity     8150 non-null    float64
4    pressure     1079 non-null    float64
5    precipintensity 3220 non-null    float64
6    cloudcover   6441 non-null    float64
7    visibility   8119 non-null    float64
8    windspeed    8020 non-null    float64
dtypes: float64(9)
memory usage: 639.8 KB

```

```
[188]: # quick check correlation
df.corr()['pm25'].sort_values()
```

```
[188]: visibility      -0.452880
temperature    -0.371182
dewpoint       -0.371018
windspeed      -0.351326
humidity       -0.082217
precipintensity  0.000103
cloudcover      0.111920
pressure        0.419824
pm25           1.000000
Name: pm25, dtype: float64
```

```
[189]: # sepearate feature and label
X = df.drop('pm25', axis=1)
y = df['pm25'].copy()
```

```
[190]: X_scaled = input_transform(data=X)
```

```
[191]: type(X_scaled)
```

```
[191]: numpy.ndarray
```

6.2 Split train and test

```
[199]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.33,
↳ random_state=2020)
```

```
[202]: len(X_train), len(y_train)
```

```
[202]: (5487, 5487)
```

```
[204]: len(X_test), len(X_test)
```

[204]: (2703, 2703)

6.3 Train and validate

```
[205]: # I will jump in and use voting (seem safer)
lin_reg = LinearRegression()
tree_reg = DecisionTreeRegressor()
rnd_reg = RandomForestRegressor()
voting_reg = VotingRegressor(
    estimators=[('lin', lin_reg),
                ('rnd', rnd_reg),
                ('tree', tree_reg)
    ],
)
```

```
[206]: voting_reg.fit(X_train, y_train)
```

```
[206]: VotingRegressor(estimators=[('lin',
                                   LinearRegression(copy_X=True, fit_intercept=True,
                                                       n_jobs=None, normalize=False)),
                                   ('rnd',
                                   RandomForestRegressor(bootstrap=True,
                                                           ccp_alpha=0.0,
                                                           criterion='mse',
                                                           max_depth=None,
                                                           max_features='auto',
                                                           max_leaf_nodes=None,
                                                           max_samples=None,
                                                           min_impurity_decrease=0.0,
                                                           min_impurity_split=None,
                                                           min_samples_leaf=1,
                                                           min_samples_split=2,
                                                           min_weight_fraction_leaf=0.0,
                                                           random_state=None, verbose=0,
                                                           warm_start=False)),
                                   ('tree',
                                   DecisionTreeRegressor(ccp_alpha=0.0,
                                                           criterion='mse',
                                                           max_depth=None,
                                                           max_features=None,
                                                           max_leaf_nodes=None,
                                                           min_impurity_decrease=0.0,
                                                           min_impurity_split=None,
                                                           min_samples_leaf=1,
                                                           min_samples_split=2,
                                                           min_weight_fraction_leaf=0.0,
```



```

presort='deprecated',
random_state=None,
splitter='best'))],
n_jobs=None, weights=None)

```

```
[207]: voting_reg.estimators_
```

```

[207]: [LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False),
RandomForestRegressor(bootstrap=True, ccp_alpha=0.0, criterion='mse',
max_depth=None, max_features='auto', max_leaf_nodes=None,
max_samples=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=100, n_jobs=None, oob_score=False,
random_state=None, verbose=0, warm_start=False),
DecisionTreeRegressor(ccp_alpha=0.0, criterion='mse', max_depth=None,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')]

```

6.3.1 Trainset

```
[208]: y_entrain_p = voting_reg.predict(X_train)
```

```

[209]: mse = mean_squared_error(y_entrain_p, y_train)
train_std = np.sqrt(mse)
train_std

```

```
[209]: 11.653075002954594
```

```
[280]: y_train.iloc[0:100]
```

```

[280]: 2018-05-17 08:00:00    20.0
2018-09-05 03:00:00    33.0
2018-12-31 10:00:00    22.0
2018-07-06 19:00:00    20.0
2018-01-05 11:00:00    88.1
...
2018-09-19 20:00:00    34.0
2018-11-15 02:00:00    20.0
2018-09-25 09:00:00    19.0
2018-09-07 05:00:00    11.0
2018-08-20 21:00:00    36.0

```

Name: pm25, Length: 100, dtype: float64

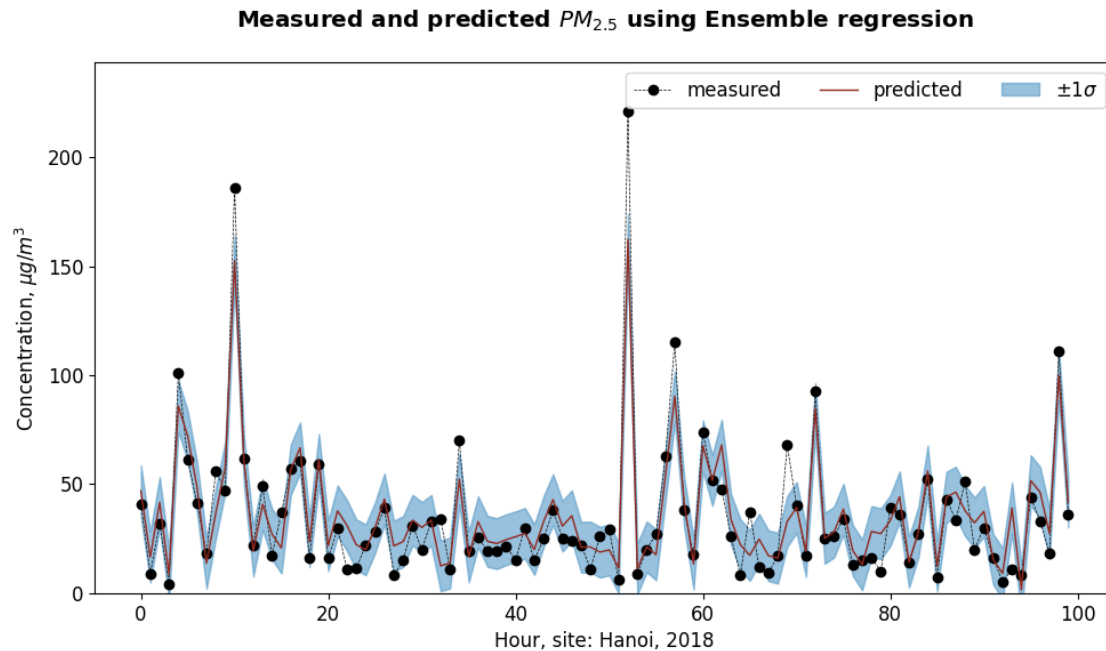
```
[211]: from random import randint
```

```
[212]: # time to invest some good graphs
def plot_results(label=None, prediction=None, std=None, points=100,
    ↪savefig=False):
    plt.style.use('default')
    plt.rcParams['font.size'] = 12
    start = randint(0, len(label)-points)
    end = start + points
    label_ = label.iloc[start:end]
    pred_ = prediction[start:end]
    xindex = np.arange(0, len(label_), 1)

    plt.figure(figsize=(10,6))
    plt.plot(label_.to_list(), 'ko--', lw=0.5, label='measured')
    plt.plot(pred_, lw=1, color='#922B21', label='predicted')
    plt.fill_between(xindex, pred_ - std_, pred_ + std_,
        color='#5499C7', alpha=0.6, label = '$\pm 1 \sigma$')

    max_ = np.max([label_.max(), np.max(pred_)])
    plt.ylim(0, 1.1*max_)
    plt.ylabel('Concentration, $\mu$ g/m3')
    plt.title('Measured and predicted $PM_{2.5}$ using Ensemble regression',
        y=1.05, weight='bold')
    plt.xlabel('Hour, site: Hanoi, 2018')
    plt.legend(ncol=3)
    if savefig:
        plt.tight_layout()
        plt.savefig(f'img/en_reg_{start}.png', optimize=True)
    return None
```

```
[213]: plot_results(label=y_train, prediction=y_entrain_p,
    std=train_std, savefig=True)
```



6.3.2 Testset

```
[214]: y_entest_p = voting_reg.predict(X_test)
```

```
[215]: mse = mean_squared_error(y_entest_p, y_test)
test_std = np.sqrt(mse)
mse, test_std
```


```
[215]: (469.53197116134754, 21.668686419839748)
```

```
[216]: add_stats(model='voting reg (Darksky)',
                train_rmse=train_std,
                test_rmse=test_std)
results
```

```
[216]: {'linear reg': {'train_rmse': 25.5, 'test_rmse': 25.9},
        'decisiontree reg': {'train_rmse': 0.0, 'test_rmse': 28.9},
        'randomforest reg': {'train_rmse': 7.3, 'test_rmse': 19.9},
        'gridsearch': {'train_rmse': 7.8, 'test_rmse': 20.0},
        'voting reg': {'train_rmse': 10.5, 'test_rmse': 21.9},
        'voting reg (Darksky)': {'train_rmse': 11.7, 'test_rmse': 21.7}}
```

- not much worse, infact, with less parameters and get a similar outcome, that is actually encouraging

```
[217]: plot_results(label=y_test,
                  prediction=y_entest_p, std=test_std, savefig=True)
```



7 RMSE

```
[218]: df = pd.DataFrame(data=results)
```

```
[219]: df
```

```
[219]:
```

	linear reg	decisiontree reg	randomforest reg	gridsearch \
train_rmse	25.5	0.0	7.3	7.8
test_rmse	25.9	28.9	19.9	20.0

	voting reg	voting reg (Darksky)
train_rmse	10.5	11.7
test_rmse	21.9	21.7

```
[222]: df2 = df.transpose()
```

```
[223]: df2
```

```
[223]:
```

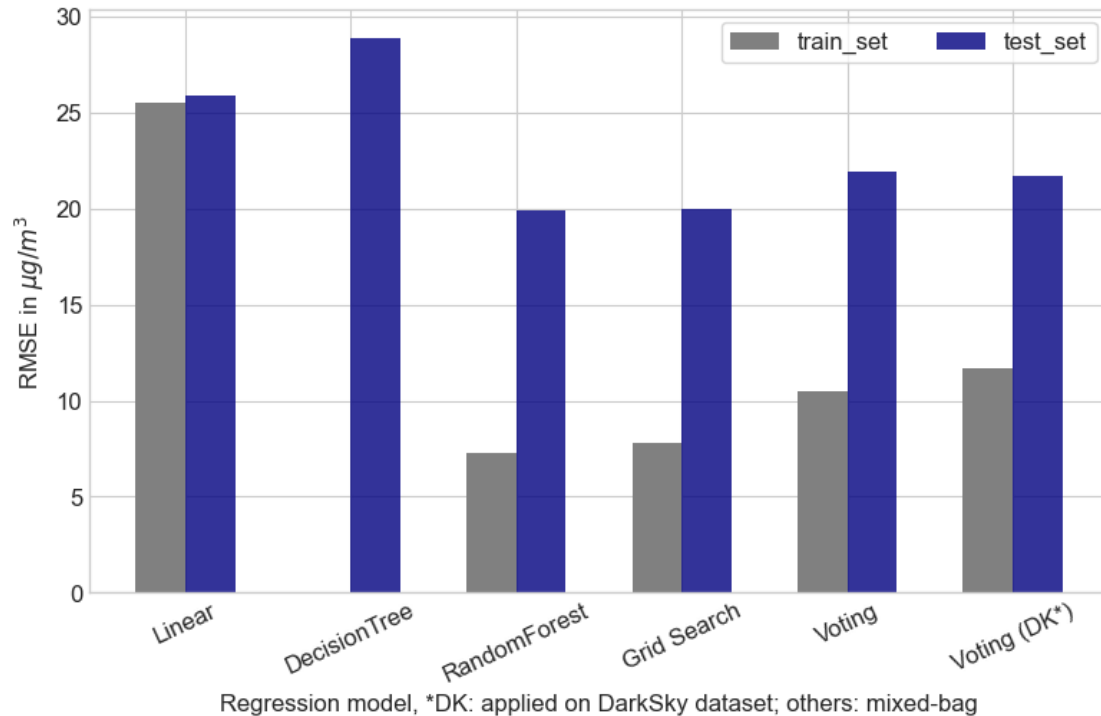
	train_rmse	test_rmse
linear reg	25.5	25.9
decisiontree reg	0.0	28.9

randomforest reg	7.3	19.9
gridsearch	7.8	20.0
voting reg	10.5	21.9
voting reg (Darksky)	11.7	21.7

```
[280]: plt.style.use('seaborn-whitegrid')
```

```
[299]: bw = 0.3
idx = np.arange(len(df2))
fig, ax = plt.subplots(figsize=(8,6))
ax.bar(idx-bw/2, df2['train_rmse'], bw, color='gray', label='train_set')
ax.bar(idx+bw/2, df2['test_rmse'], bw, color='navy', alpha=0.8,
      ↪label='test_set')
ax.set_xticklabels(['', 'Linear', 'DecisionTree', 'RandomForest',
                    'Grid Search', 'Voting', 'Voting (DK*)'],
                  rotation=25)
ax.set_xlabel('Regression model, *DK: applied on DarkSky dataset; others:
      ↪mixed-bag')
ax.set_title('Root Squared Mean Errors with $PM_{2.5}$ prediction\
      \n with meteorological parameters for Hanoi, 2018',
            y=1.05,
            weight='bold')
# labels = ax.get_xticklabels()
ax.set_ylabel('RMSE in $\mu$ g/m^3$')
ax.legend(frameon=True, ncol=2)
fig.tight_layout()
fig.savefig('img/2020Aug_rmse_raw.png')
```

Root Squared Mean Errors with $PM_{2.5}$ prediction
with meteorological parameters for Hanoi, 2018



```
[270]: pm['pm25'].mean()
```

```
[270]: 40.75225885225883
```

```
[273]: df3 = df2*100/pm['pm25'].mean()
```

```
[274]: df3
```

```
[274]:
```

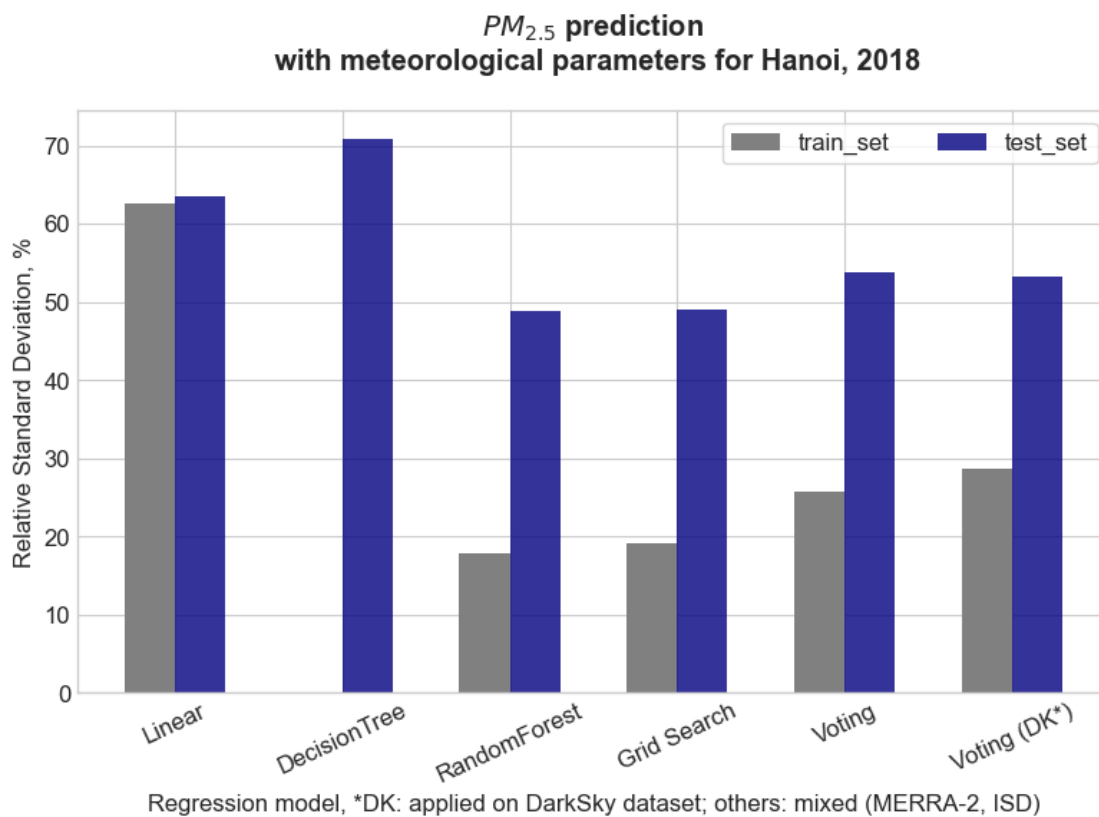
	train_rmse	test_rmse
linear reg	62.573219	63.554759
decisiontree reg	0.000000	70.916314
randomforest reg	17.913117	48.831649
gridsearch	19.140043	49.077034
voting reg	25.765443	53.739352
voting reg (Darksky)	28.710065	53.248582

```
[296]: bw = 0.3
idx = np.arange(len(df3))
fig, ax = plt.subplots(figsize=(8,6))
ax.bar(idx-bw/2, df3['train_rmse'], bw, color='gray', label='train_set')
```

```

ax.bar(idx+bw/2, df3['test_rmse'], bw, color='navy', alpha=0.8,
      label='test_set')
ax.set_xticklabels(['', 'Linear', 'DecisionTree', 'RandomForest',
                    'Grid Search', 'Voting', 'Voting (DK*)'],
                  rotation=25)
ax.set_xlabel('Regression model, *DK: applied on DarkSky dataset; others: mixed (MERRA-2, ISD)')
ax.set_title('$PM_{2.5}$ prediction\n with meteorological parameters for Hanoi, 2018',
             y=1.05,
             weight='bold')
# labels = ax.get_xticklabels()
ax.set_ylabel('Relative Standard Deviation, %')
ax.legend(frameon=True, ncol=2)
fig.tight_layout()
# fig.tight_layout(rect=(0.1,0.1,0.95, 0.95))
fig.savefig('img/2020Aug_rmse_rsd.png')

```



[]: